

Nimbus System Architecture

Nimbus uses a retrieval augmented generation pipeline. Documents are parsed, split into chunks, converted into embeddings, and stored in a vector database.

At query time, the user question is embedded and the most similar chunks are retrieved. The language model receives only the retrieved context and is instructed to answer from that context.

The reference implementation uses ChromaDB for local vector storage and the all-MiniLM-L6-v2 sentence transformer model for embeddings.